Storing Data: Disks, Buffers and Files

Begin Day 3

Berkeley cs186

Architecture of a DBMS

Organizing tables and records as groups of pages in a logical file.

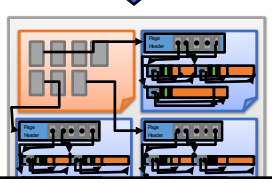

Files

Record

Bob	Harmon	M	32	94703
Varchar	Varchar	Char	Int	Int

Byte Rep. Record

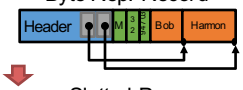

Slotted Page

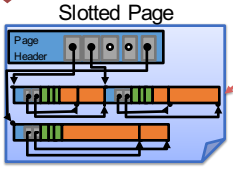

Table

File

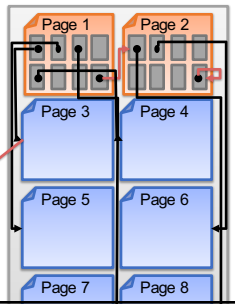

Overview

Summary

Table encoded as files which are collections of pages.

Table

Heap File

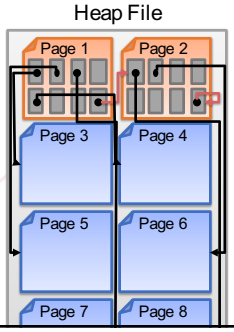

Overview

How do we store records on a page?

Table

Heap File

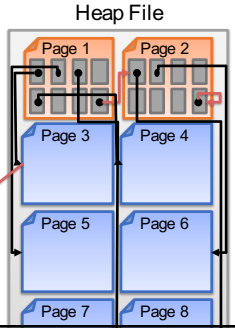

Slotted Page

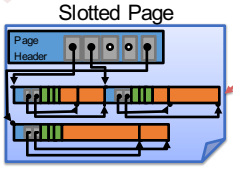

Page Basics: The Header

Blank Page

128KB

Page Basics: The Header

Header may contain:
- Number of records
- Free space
- Maybe next/last pointer
- Slot Table … more soon

Things to Address

- Record length? *Fixed* or *Variable*
- Find records by record id?
  - record id = (Page, Location in Page)
- How do we add and delete records?

Fixed Length Records: Packed

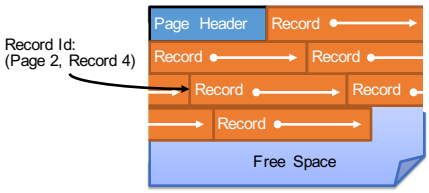

- Pack records densely
- Record id "location in page"?
  - record number in page
- Easy to add: just append
- Delete?

Fixed Length Records: Packed

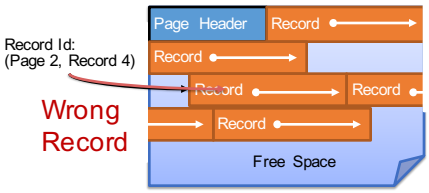

- Pack records densely
- Record id: record number in page
- Easy to add: just append
- Delete? **Packed → re-arrange** ←Necessary?

Fixed Length Records: Unpacked

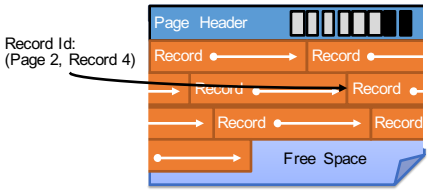

- Bitmap denotes "slots" with records
- Record id: record number in page
- **Insert:** find first empty slot
- **Delete?**

Fixed Length Records: Unpacked

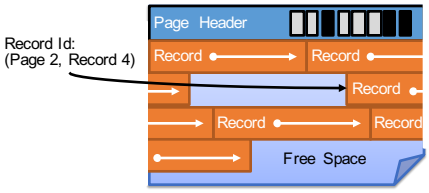

- Bitmap denotes "slots" with records
- Record id: record number in page
- **Insert:** find first empty slot
- **Delete:** Clear bit

Variable Length Records

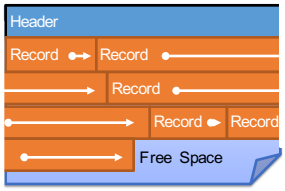

- How do we know where each record begins?
- What happens when we add and delete records?

Slotted Page

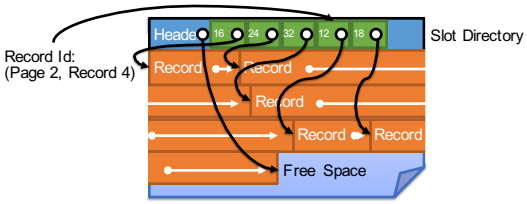

- Introduce **slot directory** in header
  - Length + Pointer to beginning of record
  - Pointer to free space
- Record ID = location in slot table
- ***Delete***? (e.g., 3rd record on the page)

Slotted Page: Delete 3rd record

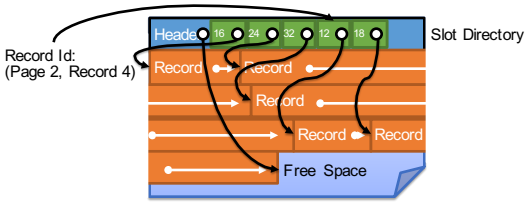

- **Delete**: Set pointer to null.
  - Doesn't affect pointers to other records
  - However, need to make sure we remove any references to **record_id** in indexes

Slotted Page

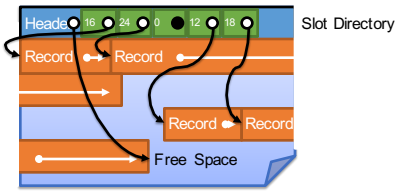

- ***Insert***? 

Slotted Page: Insert Record

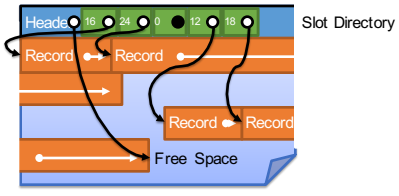

- **Insert**: 

  - Place record in free space on page

Slotted Page

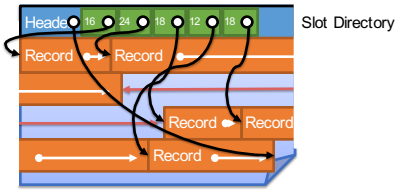

- **Insert**:
  - Place record in free space on page
  - Create pointer in next open slot in slot directory
  - Fragmentation?

Slotted Page

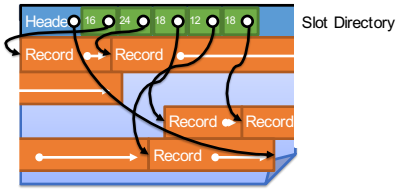

- **Insert**:
  - Place record in free space on page
  - Create pointer in next open slot in slot directory
  - Reorganize data on page.

Slotted Page

Is this safe?

When should I re-organizer?

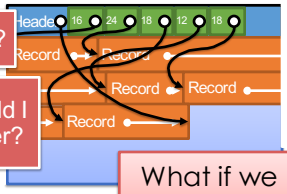

What if we need more slots?

- **Insert**:
  - Place record in free space on page
  - Create pointer in next open slot in slot directory
  - Reorganize data on page.

Slotted Page: Growing Slots

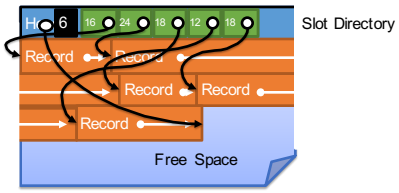

- Track number of slots in slot directory

Slotted Page: Growing Slots

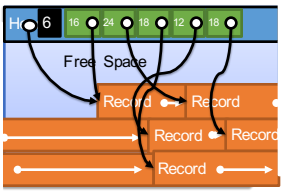

- Track number of slots in slot directory
- Grow records from other end of page
  - Why?

Slotted Page: Growing Slots

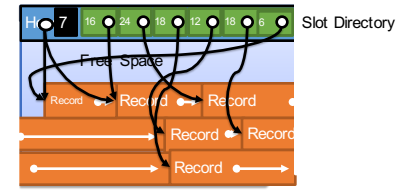

- Track number of slots in slot directory
- Grow records from other end of page
- Extend slot directory on insert
  - Add record in free space & update counter

Slotted Page: Growing Slots

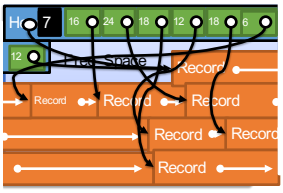

- Track number of slots in slot directory
- Grow records from other end of page
- Extend slot directory on insert
  - Add record in free space & update counter

Slotted Page Summary

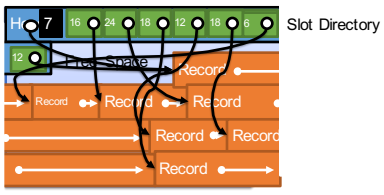

- Typically use slotted Page
- Good for fixed length records too. Why?
  - Re-arrange (e.g., sort) and squash null fields

Overview

Store records on slotted page

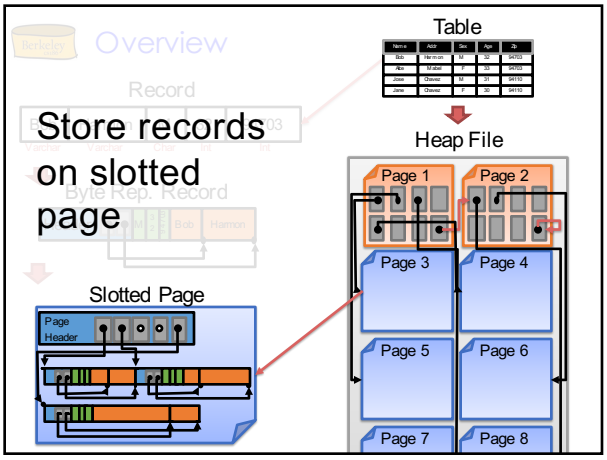

Overview

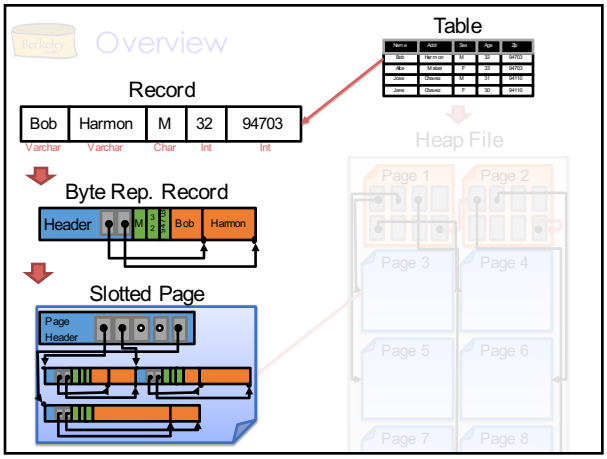

Record Formats

Record

- Relational Model →
  - Each record in table has same fixed type
- Assume *System Catalog* with *Schema*
  - No need to store type information (save space!)
  - This will be another table … (bootstraping)
- Goals:
  - Compact in memory & disk format
  - Fast access to fields (why?)

- Easy Case: *Fixed Length Fields*
- Interesting Case: *Variable Length Fields*

Record Formats: Fixed Length

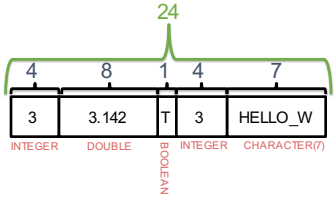

- Field types same for all records in a file.
  - Type info stored separately in *system catalog*
- On disk byte representation same as in memory
- Finding *i'th* field?
  - done via arithmetic (fast)
- Compact? (Nulls?)

Record Formats: Variable Length

What happens if fields are variable length?

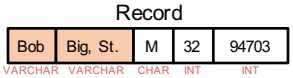

Could store with padding? (Fixed Length)

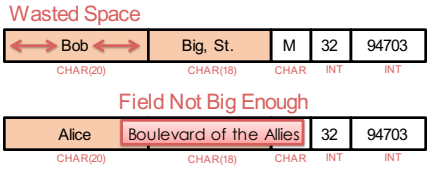

Record Formats: Variable Length

What happens if fields are variable length?

Record

Bob	Big, St.	M	32	94703
VARCHAR	VARCHAR	CHAR	INT	INT

Could use delimiters (i.e., CSV):

Comma Separated Values (CSV)

Bob	,	Big, St.	,	M	,	32	,	94703
VARCHAR		VARCHAR		CHAR		INT		INT

- Issues?

Record Formats: Variable Length

What happens if fields are variable length?

Record

Bob	Big, St.	M	32	94703
VARCHAR	VARCHAR	CHAR	INT	INT

Could use delimiters (i.e., CSV):

Comma Separated Values (CSV)

Bob	,	Big,	St.	,	M	,	32	,	94703
VARCHAR		VARCHAR			CHAR		INT		INT

- Requires scan to access field
- What if text contains commas?

Record Formats: Variable Length

What happens if fields are variable length?

Record

Bob	Big, St.	M	32	94703
VARCHAR	VARCHAR	CHAR	INT	INT

Store **length information** before fields:

Variable Length Fields with Offsets

3	Bob	8	Big, St.	M	32	94703
VARCHAR		VARCHAR		CHAR	INT	INT

- Requires scan to access field
- ~~What if text contains commas?~~

Record Formats: Variable Length

What happens if fields are variable length?

Record

Bob	Big, St.	M	32	94703
VARCHAR	VARCHAR	CHAR	INT	INT

Store **length information** before fields:

Variable Length Fields with Offsets

M	32	94703	3	Bob	8	Big, St.
CHAR	INT	INT		VARCHAR		VARCHAR

Move all variable length fields to end →enable fast access

- Requires scan to access **some** field
- ~~What if text contains commas?~~

Record Formats: Variable Length

What happens if fields are variable length?

Record

Bob	Big, St.	M	32	94703
VARCHAR	VARCHAR	CHAR	INT	INT

Introduce a **record header**:

Header	M	32	94703	Bob	Big, St.
	CHAR	INT	INT	VARCHAR	VARCHAR

- ~~Requires scan to access field.~~
- ~~What if text contains commas?~~

Record Formats: Variable Length

What happens if fields are variable length?

Record

Bob	Big, St.	M	32	94703
VARCHAR	VARCHAR	CHAR	INT	INT

Introduce a **record header**:

Header	M	32	94703	Bob	Big, St.
	CHAR	INT	INT	VARCHAR	VARCHAR

- Direct access & no "escaping", other adv.?
  - Handle null fields → useful for fixed length

Overview

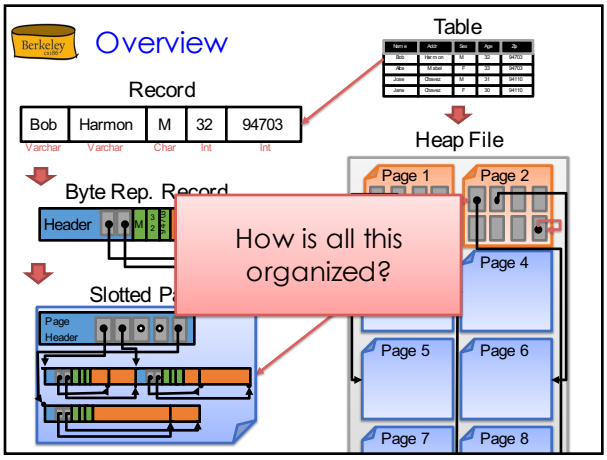

System Catalogs

- For each relation:
  - name, file location, file structure (e.g., Heap file)
  - attribute name and type, for each attribute
  - index name, for each index
  - integrity constraints
- For each index:
  - structure (e.g., B+ tree) and search key fields
- For each view:
  - view name and definition
- Plus statistics, authorization, buffer pool size, etc.

☛ *Catalogs are themselves stored as relations!*

pg_attribute

```
postgres=# \d pg_attribute
    Table "pg_catalog.pg_attribute"
    Column     |   Type    | Modifiers
---------------+-----------+-----------
 attrelid      | oid       | not null
 attname       | name      | not null
 atttypid      | oid       | not null
 attstattarget | integer   | not null
 attlen        | smallint  | not null
 attnum        | smallint  | not null
 attndims      | integer   | not null
 attcacheoff   | integer   | not null
 atttypmod     | integer   | not null
 attbyval      | boolean   | not null
 attstorage    | "char"    | not null
 attalign      | "char"    | not null
 attnotnull    | boolean   | not null
 atthasdef     | boolean   | not null
 attisdropped  | boolean   | not null
 attislocal    | boolean   | not null
 attinhcount   | integer   | not null
 attcollation  | oid       | not null
 attacl        | aclitem[] |
 attoptions    | text[]    |
 attfdwoptions | text[]    |
Indexes:
    "pg_attribute_relid_attnam_index" UNIQUE, btree (attrelid, attname)
    "pg_attribute_relid_attnum_index" UNIQUE, btree (attrelid, attnum)

postgres@precise64:~$
```

sqlite_master

```
SELECT name, rootpage FROM sqlite_master
WHERE type='table'
ORDER BY name;
```

```
sqlite>
sqlite>
sqlite>
sqlite>
sqlite> SELECT name, rootpage FROM sqlite_master
   ...> WHERE type='table'
   ...> ORDER BY name;
name         rootpage
-----------  ----------
Album        2
Artist       3
Customer     4
Employee     7
Genre        9
Invoice      10
InvoiceLin   12
MediaType    14
Playlist     15
PlaylistTr   16
Track        19
sqlite>
```

Summary

- Disk manager loads and stores pages
  - Block level reasoning
  - Abstracts device and file system; provides fast next
- Buffer manager brings pages into RAM
  - page pinned while reading/writing
  - dirty pages written to disk
  - good *replacement policy* essential for performance
- DBMS "File" tracks collection of pages, records within each.
  - Heap-files: unordered records organized with directories

Summary (Contd.)

- Slotted page format
  - Variable length records and intra-page reorg
- Variable length record format
  - Direct access to i'th field and null values.
- Catalog relations store information about relations, indexes and views.

Architecture of a DBMS

Organized in layers

Each layer abstracts the layer below
- Manage complexity
- Perf. Assumptions

Example of **good** systems design

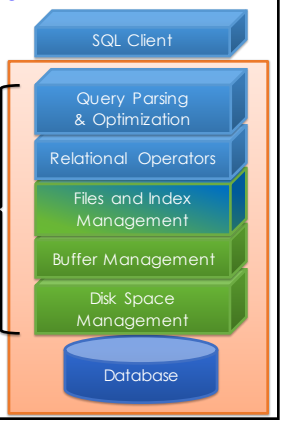

File Organizations and Indexing

Berkeley cs186

"If you don't find it in the index, look very carefully through the entire catalogue."
— Sears, Roebuck, and Co., Consumer's Guide, 1897

Architecture of a DBMS

Organized in layers

Each layer abstracts the layer below
- Manage complexity
- Perf. Assumptions

Example of **good** systems design

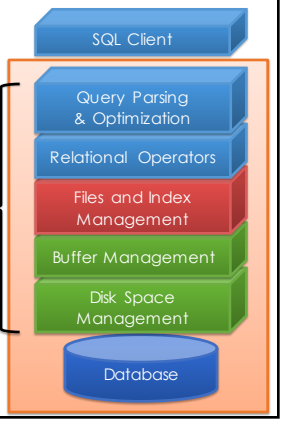

Last Time:

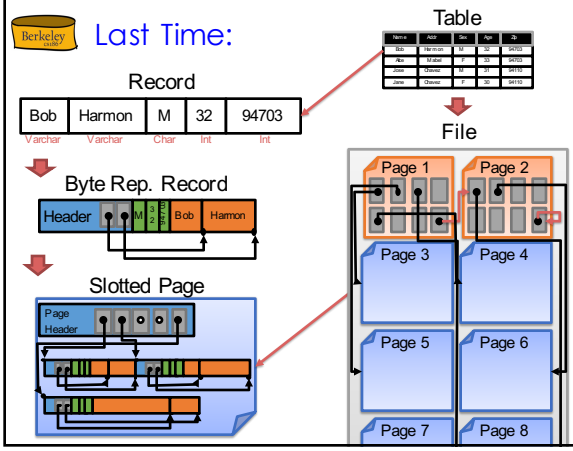

How do we organize records and pages for certain operations?

- Ordering and grouping records

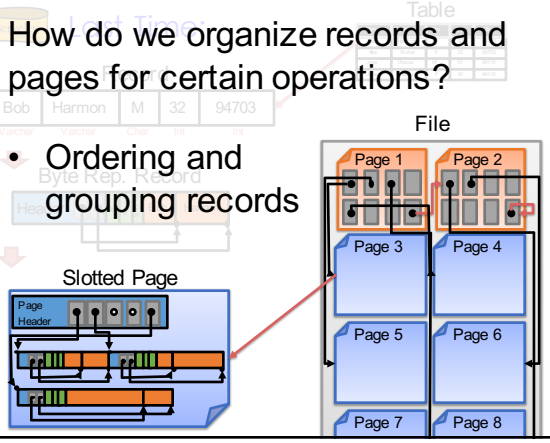

Last Time Discussed Heap Files

Heap Files
- Unordered collection of records
- Add/Remove records: Easy (Cost?)
- Scan: Easy (Cost?)
- Find a record?
  - Given a record_id: (File, Page, Slot)?
  - Matching username = "joeyg"?

Multiple File Organizations

Many alternatives exist, *each good in some situations and not so good in others*:

- Heap files: Suitable when typical access is a full scan of all records
- Sorted Files: Best for retrieval in *search key* order, or a **range** of records is needed
- Clustered Files & Indexes: Group data into blocks to enable fast lookup and efficient modification. (More on this soon …)

Bigger Questions

- What is the "best" file organization?
  - Depends on access patterns…
  - how? what are they?
- Can we be quantitative about tradeoffs?
  - Better → How Much?

Goals for Today

- Big picture overheads for data access
  - We'll (overly) simplify things to **gain insight**
  - Still, a bit of discipline:
    - Clearly identify assumptions up front
    - Then estimate cost in a principled way
- Foundation for query optimization
  - Can't choose the fastest scheme without an estimate of speed!

Cost Model for Analysis

- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** (Average) time to read or write disk block
- *Average-case* analyses for *uniform random* workloads
- We will ignore:
  - Sequential vs. Random I/O
  - Pre-fetching
  - Any in-memory costs

☛ ***Good enough to show the overall trends!***

More Assumptions

- Single record insert and delete.
- Equality selection - exactly one match
- For Heap Files:
  - Insert always appends to end of file.
- For Sorted Files:
  - Files compacted after deletions.
  - Sorted according to search key.
- Question all these assumptions and rework
  - As an exercise to study for tests, generate ideas

Heap Files & Sorted Files

B: The number of data pages = 5
R: Number of records per page = 2
D: (Average) time to read or write disk page = 5 ms

Heap File

Sorted File

Records are just integers

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records		
Equality Search		
Range Search		
Insert		
Delete		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records		
Equality Search		
Range Search		
Insert		
Delete		

Scan all the Records?

Heap File

2, 5 | 1, 6 | 4, 7 | 3, 10 | 8, 9

Sorted File

1, 2 | 3, 4 | 5, 6 | 7, 8 | 9, 10

B: data pages
R: records per page
D: Avg. IO Time

Pages touched: ?
Time to read the record: ?

http://hyperboleandahalf.blogspot.com

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search		
Range Search		
Insert		
Delete		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search		
Range Search		
Insert		
Delete		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search		
Range Search		
Insert		
Delete		

Find Key 8

B: data pages
R: records per page
D: Avg. IO Time

Heap File

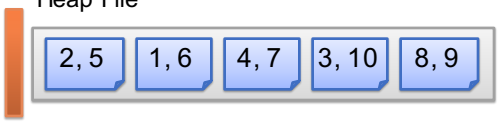

Pages touched **on average**?

- **P**(i): Probability of key on page **i** is: **1/B**
- **T**(i): # Pages touched if key on page **i** is: **i**
- Therefore the expected # pages touched:

$$\sum_{i=1}^{B} \mathrm{T}(i)\mathbf{P}(i) = \sum_{i=1}^{B} i\frac{1}{B} = \frac{B(B+1)}{2B} \approx \frac{B}{2}$$

Find Key 8

B: data pages
R: records per page
D: Avg. IO Time

Heap File

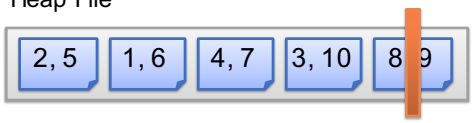

Pages touched **on average**: B/2

Breaking an Assumption:
What if there was more than one key?

- Need to check all pages → B

Find Key 8

B: data pages
R: records per page
D: Avg. IO Time

Sorted File

Binary Search

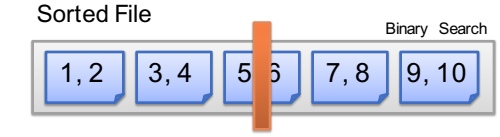

Worst-case: Pages touched in binary search

- **Log_2B**

Average-case: Pages touched in binary search

- **Log_2B?**

Average Case Binary Search

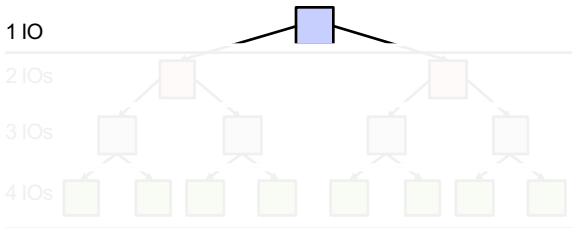

Expected Number of Reads: 1 (1 / B) + 2 (2 / B) + 3 (4 / B) + 4 (8 / B)

$$\sum_{i=1}^{\log_2 B} i\frac{2^{i-1}}{B} = \frac{1}{B}\sum_{i=1}^{\log_2 B} i2^{i-1} = \log_2 B - \frac{B-1}{B}$$

Average ≈ Worst

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	($\log_2$ B) * D
Range Search		
Insert		
Delete		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	($\log_2$ B) * D
Range Search		
Insert		
Delete		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search		
Insert		
Delete		

Keys between 7 and 9

B: data pages
R: records per page
D: Avg. IO Time

Heap File

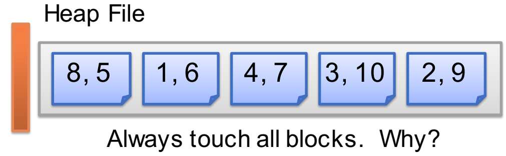

Always touch all blocks. Why?

Keys between 7 and 9

B: data pages
R: records per page
D: Avg. IO Time

Heap File

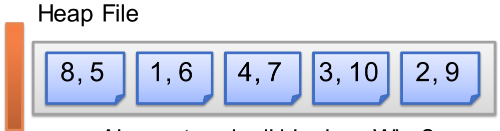

Always touch all blocks. Why?

Sorted File

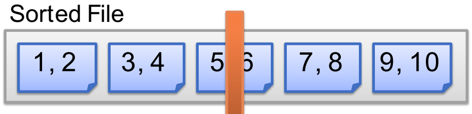

1. Find beginning of range
2. Scan right

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$
Insert		
Delete		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$
Insert		
Delete		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$
Insert		
Delete		

Insert 4.5

B: data pages
R: records per page
D: Avg. IO Time

Heap File

2, 5 | 1, 6 | 4, 7 | 3, 10 | 8, 9 | 4.5,

Stick at the end of the file. Cost? ***2D***

Why 2?

Insert 4.5

B: data pages
R: records per page
D: Avg. IO Time

Heap File

2, 5 | 1, 6 | 4, 7 | 3, 10 | 8, 9 | 4.5,

Read last page, append, write. ***2D***

Sorted File

1, 2 | 3, 4 | 5, 6 | 7, 8 | 9, 10

1. Find location for record: **Log_2B**

Insert 4.5

B: data pages
R: records per page
D: Avg. IO Time

Heap File

2, 5 | 1, 6 | 4, 7 | 3, 10 | 8, 9 | 4.5,

Read last page, append, write. ***2D***

Sorted File

1, 2 | 3, 4 | 4.5,5 | 6, 7 | 8, 9 | 10,

1. Find location for record: **Log_2B**
2. Insert and shift rest of file Cost? **2 (B/2) Why?**

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$
Insert	2D	$((\log_2 B)+B)D$
Delete		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$
Insert	2D	$((\log_2 B)+B)D$
Delete		

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$
Insert	2D	$((\log_2 B)+B)D$
Delete		

Delete 4.5

B: data pages
R: records per page
D: Avg. IO Time

Heap File

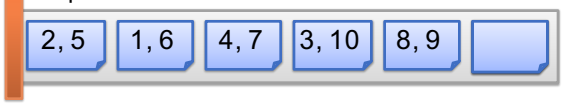

Average case to find record: **B/2 Reads**
Delete record from page.

Cost? (B/2 + 1)D

Why + 1?

Delete 4.5

B: data pages
R: records per page
D: Avg. IO Time

Heap File

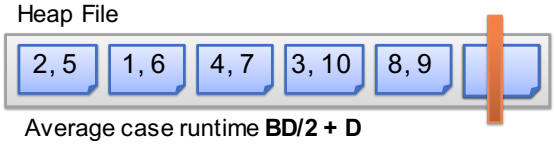

Average case runtime **BD/2 + D**

Sorted File

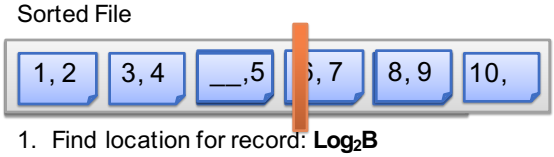

1. Find location for record: **Log_2B**
2. Delete record in page → Gap!

Delete 4.5

B: data pages
R: records per page
D: Avg. IO Time

Heap File

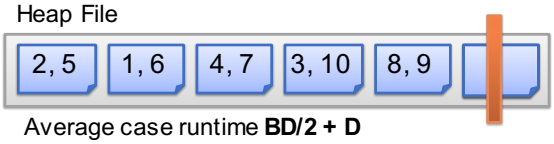

Average case runtime **BD/2 + D**

Sorted File

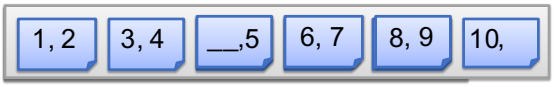

1. Find location for record: **Log_2B**
2. Shift rest of file left by 1 record: **2 (B/2)**

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$
Insert	2D	$((\log_2 B)+B)D$
Delete	0.5BD + D	$((\log_2 B)+B)D$

Which is better?

Day 2: File Organizations and Indexing

Berkeley cs186

"If you don't find it in the index, look very carefully through the entire catalogue."
— Sears, Roebuck, and Co., Consumer's Guide, 1897

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$
Insert	2D	$((\log_2 B)+B)D$
Delete	0.5BD + D	$((\log_2 B)+B)D$

Issues:
- Find
- Range
- Modification

Can we do better?

Indexes

Overview: Indexes

*An **index** is a data structure that enables fast lookup of data entries by search key*

- **Lookup:** may support many different operations
  - Equivalence, range, region, …
- **Search Key:** any *subset* of columns in the relation
  - Do not need to be unique
- **Data Entries:** items stored in the index, (**k**, {items})
  - Could be actual records or record ids (pointers)…
- **Many Types:** B-Tree, Hash, ISAM, R-Tree, GiST, …
- **Modification:** often support fast insert and delete
  - Lazily maintain ordering, clustering, …

Kinds of Lookups Supported?

- Basic Selection: <key> <op> <constant>
  - Equality selections (op is =)?
  - Range selections (op is one of <, >, <=, >=, BETWEEN)
- More exotic selections:
  - 2-dimensional ranges ("east of Berkeley and west of Truckee and North of Fresno and South of Eureka")
  - 2-dimensional radii ("within 2 miles of Soda Hall")
  - Common **n-dimensional index**: *R-tree, KD-Tree*
    - Beware of the ***curse of dimensionality***
  - Ranking queries ("10 restaurants closest to Berkeley")
  - Regular expression matches, genome string matches, etc.
  - See http://en.wikipedia.org/wiki/GiST for more

Search Key: *Any* Subset of Columns

- Search key needn't be a key of the relation
  - Recall: key of a relation must be unique (e.g., SSN)
  - Search keys don't have to be unique
- **Composite Keys**: more than one column
  - Think: Phone Book <Last Name, First>
  - **Lexicographic order**
  - <Age, Salary>:
    - ✓ Age = 31 & Sal = 400
    - ✓ Age = 55 & Sal > 200
    - ✗ Age > 31 & Sal = 400
    - ✓ Age = 31
    - ✓ Age > 31
    - ✗ Sal = 300

SSN	Last Name	First Name	Age	Salary
123	Adams	Elmo	31	$400
443	Grouch	Oscar	32	$300
244	Oz	Bert	55	$140
134	Sanders	Ernie	55	$400

✗ Means that the index is unable able to exclude all entries that are not in the result set.

Data Entries: How are they stored?

- What is the representation of data in the index?
  - Actual data or pointer(s) to the data
- How is the data stored in the data file?
  - Clustered or unclusted with respect to the indexed
- Big Impact on Performance

How is data stored in the index

- Three alternatives:
  1. **By Value:** actual data record (with key value **k**)
  2. **By Reference:** <**k**, rid of matching data record>
  3. **By List of Refs.:** <**k**, list of rids of *all* matching data records>
- Choice is orthogonal to the indexing technique.
  - B+ trees, hash-based structures, R trees, GiSTs, …
- Can have multiple (different) indexes per file.
  - E.g. file sorted by *age*, with a hash index on *salary* and a B+tree index on *name*.

Alternatives for Data Entries (Contd.)

Alternative 1 (By Value):
Actual data record (with key value ***k***)

- Index as a file organization for records
  - Alongside heap files or sorted files
- No "**pointer lookups**" to get data records
  - Following record ids
- Could a single relation have multiple indexes of this form?
  - Probably but it would be a bad idea. Why?

Alternatives for Data Entries (Contd.)

Alternative 2 (By Reference)
<**k**, rid of matching data record>
and Alternative 3 (By List of References)
<**k**, list of rids of matching data records>

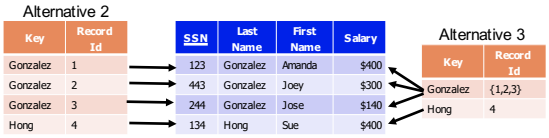

- Alts. 2 or 3 needed to support multiple indexes per relation!
- Alt. 3 more compact than Alt. 2
- For very large rid lists, single data entry spans multiple blocks

Clustered vs. Unclustered Index

- In a clustered index:
  - index data entries are stored in (approximate) order by value of search keys in data records

Clustered vs. Unclustered Index

- In a clustered index:
  - index data entries are stored in (approximate) order by value of search keys in data records

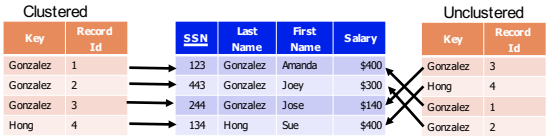

Clustered vs. Unclustered Index

- In a clustered index:
  - index data entries are stored in (approximate) order by value of search keys in data records
  - A file can be clustered on at most one search key.
  - Cost of retrieving data records through index varies *greatly* based on whether index is clustered or not!
  - Can Alternative 1 be un-clustered?
    - No, but clustered does not imply alt. 1 (earlier figure)
- Note: there is another definition of "clustering"
  - **Data Mining/AI**: grouping similar items in n-space

Clustered vs. Unclustered Index

Alternative 2 (by ref) data entries, data records in a Heap file.

- To build clustered index, first sort the heap file
  - Leave some free space on each block for future inserts
- Overflow blocks may be needed for inserts.
  - Thus, order of data recs is "close to", but not identical to, the sort order.

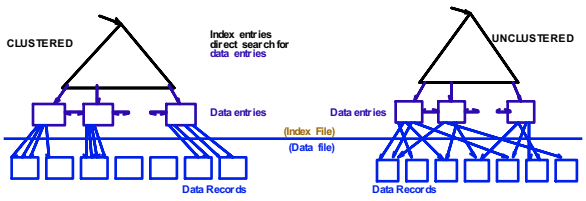

Clustered vs. Unclustered Index

Alternative 2 (by ref) data entries, data records in a Heap file.

- To build clustered index, first sort the Heap file
  - Leave some free space on each block for future inserts
- Overflow blocks may be needed for inserts.
  - Thus, order of data recs is "close to", but not identical to, the sort order.

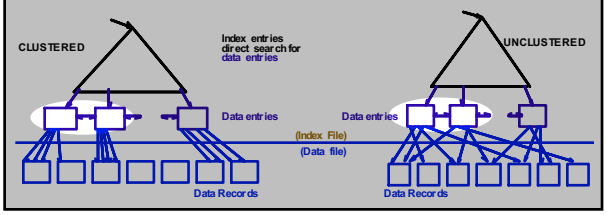

Clustered vs. Unclustered Index

Alternative 2 (by ref) data entries, data records in a Heap file.

- To build clustered index, first sort the Heap file
  - Leave some free space on each block for future inserts
- Overflow blocks may be needed for inserts.
  - Thus, order of data recs is "close to", but not identical to, the sort order.

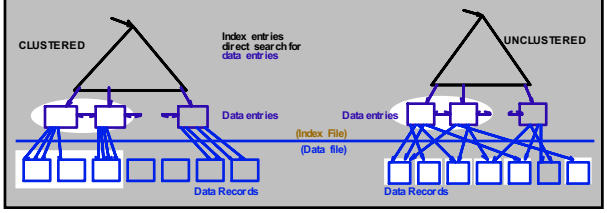

Unclustered vs. Clustered Indexes

- Clustered Pros
  - Efficient for range searches
  - Potential locality benefits?
    - Sequential disk access, prefetching, etc.
  - Support certain types of compression
    - More soon on this topic
- Clustered Cons
  - More expensive to maintain
    - Need to update index data structure
    - Solution: on the fly or "lazily" via reorgs
  - Heap file usually only **packed to 2/3** to accommodate inserts

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File
Scan all records	BD	BD
Equality Search	0.5 BD	$(\log_2 B) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$
Insert	2D	$((\log_2 B)+B)D$
Delete	0.5BD + D	$((\log_2 B)+B)D$

Issues:
- Find
- Range
- Modification

Can we do better?

Indexes

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	
Equality Search	0.5 BD	$(\log_2 B) * D$	
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$	
Insert	2D	$((\log_2 B)+B)D$	
Delete	0.5BD + D	$((\log_2 B)+B)D$	

Simple Clustered Index Analysis

Assumptions:

- Store data by reference (Alternative 2)
- Clustered tree index with 2/3 full heap file pages
  - Clustered → Heap file is initially sorted
  - **Fan-out** (F): relatively large. Why?
    - Page of <key, pointer> pairs ~ O(R)
  - Assume static index

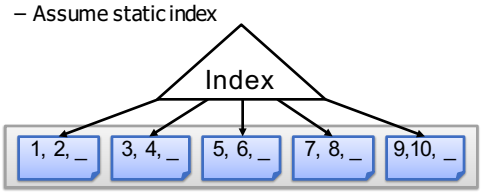

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	?
Equality Search	0.5 BD	$(\log_2 B) * D$	
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$	
Insert	2D	$((\log_2 B)+B)D$	
Delete	0.5BD + D	$((\log_2 B)+B)D$	

Scan all the Records?

B: The number of data pages **(originally)**
R: Number of records per page **(originally)**
D: (Average) time to read or write disk page

Assumptions:

- Store data by reference (Alternative 2)
- Clustered tree index with 2/3 full heap file pages
  - **Clustered** → Heap file is initially sorted
  - **Fan-out** (F): relatively large. ~ O(R)
  - Assume static index

Do we need the index? *No*

Cost? = 1.5 BD Why? = (3/2) BD

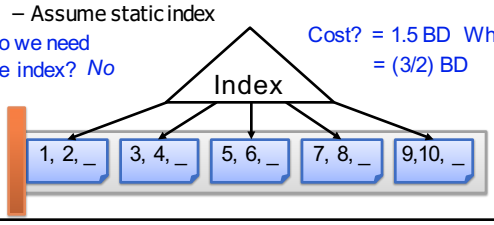

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	1.5BD
Equality Search	0.5 BD	($\log_2$ B) * D	
Range Search	BD	[($\log_2$ B) + #match pg]*D	
Insert	2D	(($\log_2$B)+B)D	
Delete	0.5BD + D	(($\log_2$B)+B)D	

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	1.5BD
Equality Search	0.5 BD	($\log_2$ B) * D	?
Range Search	BD	[($\log_2$ B) + #match pg]*D	
Insert	2D	(($\log_2$B)+B)D	
Delete	0.5BD + D	(($\log_2$B)+B)D	

Find the record with key 3

Search the index:

- Each page load narrows search by **factor of F**

Cost? = Log_F(1.5 B) D

Lookup record in heap file by **record-id**

- Recall record-id = <page, slot #>

Cost? =D

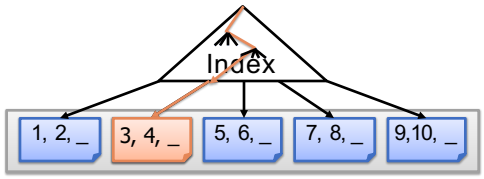

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	(3/2) BD = 1.5BD
Equality Search	0.5 BD	($\log_2$ B) * D	($\log_F$ 1.5B+1) * D
Range Search	BD	[($\log_2$ B) + #match pg]*D	
Insert	2D	(($\log_2$B)+B)D	
Delete	0.5BD + D	(($\log_2$B)+B)D	

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	(3/2) BD = 1.5BD
Equality Search	0.5 BD	($\log_2$ B) * D	($\log_F$ 1.5B+1) * D
Range Search	BD	[($\log_2$ B) + #match pg]*D	?
Insert	2D	(($\log_2$B)+B)D	
Delete	0.5BD + D	(($\log_2$B)+B)D	

Find keys between 3 and 7

Search the index for 3:= $Log_F(1.5\,B)\,D$
- Each page load narrows search by **factor of F**

Lookup record in heap file by **record-id**= D
- Recall record-id = <page, slot #>

Scan the leaves of index until the end of range (assume good clustering) = (#match pg) D

Cost?

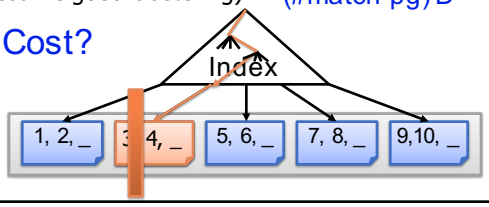

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	(3/2) BD = 1.5BD
Equality Search	0.5 BD	$(\log_2 B) * D$	$(\log_F 1.5B+1) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg]*D	$[(\log_F 1.5B) +$ #match pg]*D
Insert	2D	$((\log_2 B)+B)D$	?
Delete	0.5BD + D	$((\log_2 B)+B)D$	

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	(3/2) BD = 1.5BD
Equality Search	0.5 BD	$(\log_2 B) * D$	$(\log_F 1.5B+1) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg]*D	$[(\log_F 1.5B) +$ #match pg]*D
Insert	2D	$((\log_2 B)+B)D$	$((\log_F 1.5B)+2)*D$
Delete	0.5BD + D	$((\log_2 B)+B)D$	

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	(3/2) BD = 1.5BD
Equality Search	0.5 BD	$(\log_2 B) * D$	$(\log_F 1.5B+1) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg]*D	$[(\log_F 1.5B) +$ #match pg]*D
Insert	2D	$((\log_2 B)+B)D$	$((\log_F 1.5B)+2)*D$
Delete	0.5BD + D	$((\log_2 B)+B)D$	$((\log_F 1.5B)+2)*D$

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	(3/2) BD = 1.5BD
Equality Search	0.5 BD	$(\log_2 B) * D$	$(\log_F 1.5B+1) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg]*D	$[(\log_F 1.5B) +$ #match pg]*D
Insert	2D	$((\log_2 B)+B)D$	$((\log_F 1.5B)+2)*D$
Delete	0.5BD + D	$((\log_2 B)+B)D$	$((\log_F 1.5B)+2)*D$

Cost of Operations

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

	Heap File	Sorted File	Clustered Index
Scan all records	BD	BD	(3/2) BD = 1.5BD
Equality Search	0.5 BD	$(\log_2 B) * D$	$(\log_F 1.5B+1) * D$
Range Search	BD	$[(\log_2 B) +$ #match pg$]*D$	$[(\log_F 1.5B) +$ #match pg$]*D$
Insert	2D	$((\log_2 B)+B)D$	$((\log_F 1.5B)+2)*D$
Delete	0.5BD + D	$((\log_2 B)+B)D$	$((\log_F 1.5B)+2)*D$

Summary

- Many file orgs, with tradeoffs
  - Heap Files, Sorted Files, Clustered Files and Indexes
  - Benefits depend on the common operations
  - Compute expected costs
- Indexes: fast lookup of data entries by search key
  - Lookup: equivalence, range, region …
  - Search key: arbitrary columns
- Data entries:
  - 3 alternatives: By Value, By Reference, By List of References
- Often multiple indexes per file of data records
  - each with a different search key.
- Indexes can be classified as *clustered* vs. *unclustered*
  - Difs have important consequences for utility/performance.

Architecture of a DBMS

Today:

- Cost of access patterns for different **file formats**
- Introduced **indexes** as a mechanism to reduce costs

Thursday:

- We go deep on **tree indexes**: *B+-Trees*