Crash Recovery

R&G - Chapter 20

Review: The ACID properties

- Atomicity: All actions in the Xact happen, or none happen.
- Consistency: If each Xact is consistent, and the DB starts consistent, it ends up consistent.
- Isolation: Execution of one Xact is isolated from that of other Xacts.
- Durability: If a Xact commits, its effects persist.

- Question: which ones does the Recovery Manager help with?

  Atomicity & Durability (and also used for Consistency-related rollbacks)

Motivation

- **Atomicity:**
  - Transactions may abort ("Rollback").
- **Durability:**
  - What if DBMS stops running? (Causes?)

- Desired state after system restarts:
  - T1 & T3 should be durable.
  - T2, T4 & T5 should be aborted (effects not seen).

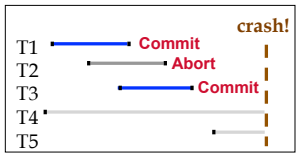

Assumptions

- **Concurrency control is in effect.**
  - Strict 2PL, in particular.
- **Updates are happening "in place".**
  - i.e. data is overwritten on (deleted from) the actual page copies (not private copies).

- **Can you think of a simple scheme (requiring no logging) to guarantee Atomicity & Durability?**
  - What happens during normal execution
  - What happens when a transaction commits?
  - What happens when a transaction aborts?

Buffer Mgmt Plays a Key Role

- **Force policy – make sure that every update is on the DB disk before commit.**
  - Provides durability without REDO logging.
  - But, can cause poor performance.

- **No Steal policy – don't allow buffer-pool frames with uncommited updates to overwrite committed data on DB disk.**
  - Useful for ensuring atomicity without UNDO logging.
  - But can cause poor performance.

In practice, even to get Force/NoSteal to work requires some nasty details for handling unexpected failures…

Preferred Policy: Steal/No-Force

- Most complicated, but highest performance.
- NO FORCE **(complicates enforcing Durability)**
  - What if system crashes before a modified page written by a committed transaction makes it to DB disk?
    - Write as little as possible, in a convenient place, at commit time, to support REDOing modifications.
- STEAL **(complicates enforcing Atomicity)**
  - What if a Xact that performed updates aborts?
  - What if system crashes before Xact is finished?
    - Must remember the old value of P (to support UNDOing the write to page P).

Buffer Management summary

	No Steal	Steal
No Force		Fastest
Force	Slowest	

Performance Implications

	No Steal	Steal
No Force	No UNDO REDO	UNDO REDO
Force	No UNDO No REDO	UNDO No REDO

Log/Recovery Implications

Basic Idea: Logging

- **Record REDO and UNDO information, for every update, in a *log*.**
  - Sequential writes to log (on a separate disk).
  - Minimal info (diff) written to log
    - Multiple updates fit in a single log page!
- **Log: An ordered list of REDO/UNDO actions**
  - Log record contains:
    <XID, pageID, offset, length, old data, new data>
  - and additional control info (which we'll see soon).

Write-Ahead Logging (WAL)

- **The Write-Ahead Logging Protocol:**
  ① Must force the log record for an update *before* the corresponding data page gets to disk.
  ② Must force all log records for a Xact *before commit*.

  I.e. transaction is not committed until all of its log records—including its "commit" record—are on the stable log.

- **#1 (with UNDO info) helps guarantee Atomicity.**
- **#2 (with REDO info) helps guarantee Durability.**
- **This allows us to implement Steal/No-Force**

- **Exactly how is logging (and recovery!) done?**
  - We'll look at the ARIES algorithms from IBM.

WAL & the Log

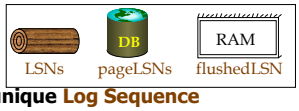

- **Each log record has a unique Log Sequence Number (LSN).**
  - LSNs always increasing.
- **Each *data page* contains a pageLSN.**
  - The LSN of the most recent *log record* for an update to that page.
- **System keeps track of flushedLSN.**
  - The max LSN flushed so far.
- **WAL: Before page i is written to DB log must satisfy:**
  $pageLSN_i \le flushedLSN$

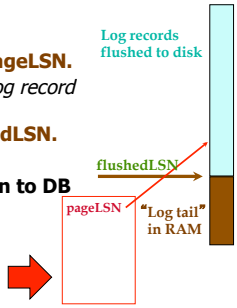

Log Records

LogRecord fields:

LSN
prevLSN
XID
type
update records only: pageID, length, offset, before-image, after-image

prevLSN is the LSN of the previous log record written by *this* Xact
(records of an Xact form a linked list backwards in time)

Possible log record types:

- Update, Commit, Abort
- Checkpoint (for log maintainence)
- Compensation Log Records (CLRs)
  - for UNDO actions
- End **(end of commit or abort)**

Other Log-Related State

- **Two in-memory tables:**
- **Transaction Table**
  - One entry per currently active Xact.
    - entry removed when Xact commits or aborts
  - Contains XID, status (running/committing/aborting), and lastLSN (most recent LSN written by Xact).
- **Dirty Page Table:**
  - One entry per dirty page currently in buffer pool.
  - Contains recLSN -- the LSN of the log record which ***first*** caused the page to be dirty.

The Big Picture: What's Stored Where

LogRecords
- LSN
- prevLSN
- XID
- type
- pageID
- length
- offset
- before-image
- after-image

Data pages each with a pageLSN

Master record

Xact Table
- lastLSN
- status

Dirty Page Table
- recLSN

flushedLSN

Buffer pool
Log tail

Normal Execution of an Xact

- **Series of reads & writes, followed by commit or abort.**
  - We will assume that disk write is atomic.
    - In practice, additional details to deal with non-atomic writes.
- **Strict 2PL.**
- **STEAL, NO-FORCE buffer management, with Write-Ahead Logging.**

Transaction Commit

- **Write commit record to log.**
- **All log records up to Xact's commit record are flushed to disk.**
  - Guarantees that flushedLSN ≥ lastLSN.
  - Note: log flushes are sequential writes to disk.
    - Can happen in asynchronous batches for efficiency.
  - Many log records per log page.
- **Commit() returns.**
- **Write end record to log.**

Simple Transaction Abort

- **For now, consider an explicit abort of a Xact.**
  - No crash involved.
- **We want to "play back" the log in reverse order, UNDOing updates.**
  - Get lastLSN of Xact from Xact table.
  - Write a new *Abort* log record to the end of the log before starting to rollback operations
  - Can follow chain of log records backward via the prevLSN field.
  - Write a "CLR" (compensation log record) for each undone operation.

Abort, cont.

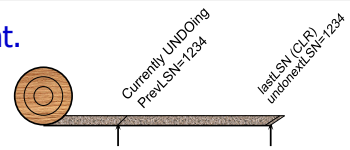

- **To perform UNDO, must have a lock on data!**
  - No problem! (why?)
- **Before restoring old value of a page, write a CLR @end of log:**
  - You continue logging while you UNDO!!
  - CLR has one extra field: undonextLSN
    - Points to the next LSN to undo (i.e. the prevLSN of the record we're currently undoing).
  - CLR contains REDO info
  - CLRs *never* Undone
    - Undo needn't be idempotent (>1 UNDO won't happen)
    - But they might be Redone when repeating history (=1 UNDO guaranteed)
- **At end of all UNDOs, write an "end" log record.**

Checkpointing

- **Conceptually, keep log around for all time.**
  - Performance/implementation problems…
- **Periodically, the DBMS creates a checkpoint**
  - Minimizes recovery time after crash. Write to log:
    - begin_checkpoint record: Indicates when chkpt began.
    - end_checkpoint record: Contains current *Xact table* and *dirty page table*. A 'fuzzy checkpoint':
      - Other Xacts continue to run; so these tables accurate only as of the time of the begin_checkpoint record.
      - No attempt to force dirty pages to disk; effectiveness of checkpoint limited by oldest unwritten change to a dirty page.
    - Store LSN of most recent chkpt record in a safe place (*master* record).

Crash Recovery: Big Picture

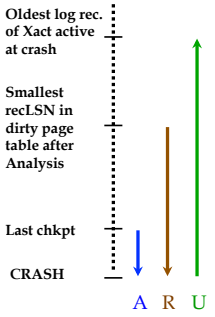

- Start from a checkpoint (found via master record).
- Three phases. Need to do:
  - Analysis - Figure out which Xacts committed since checkpoint, which failed.
  - REDO *all* actions. (repeat history)
  - UNDO effects of failed Xacts.

Recovery: The Analysis Phase

- **Re-establish knowledge of state at checkpoint.**
  - via transaction table and dirty page table stored in the checkpoint
- **Scan log forward from checkpoint.**
  - End record: Remove Xact from Xact table.
  - All Other records: Add Xact to Xact table, set lastLSN=LSN, change Xact status on commit.
  - also, for Update records: If page P not in Dirty Page Table, Add P to DPT, set its recLSN=LSN.
- **At end of Analysis…**
  - Xact table says which xacts were active ~~at time of crash.~~ at last log flush before crash
  - DPT says which dirty pages *might not* have made it to disk

Phase 2: The REDO Phase

- **We *Repeat History* to reconstruct state at crash:**
  - Reapply *all* updates (even of aborted Xacts!), redo CLRs.
- **Scan forward from log rec containing smallest recLSN in DPT.** Q: why start here?
- **For each update log record or CLR with a given LSN, REDO the action unless:**
  - Affected page is not in the Dirty Page Table, or
  - Affected page is in D.P.T., but has recLSN > LSN, or
  - pageLSN (in DB) ≥ LSN. (this last case requires I/O)
- **To REDO an action:**
  - Reapply logged action.
  - Set pageLSN to LSN. No additional logging, no forcing!

Phase 3: The UNDO Phase

- **A Naïve solution:**
  - The xacts in the Xact Table are losers.
  - For each loser, perform simple transaction abort.
  - Problems?

Phase 3: The UNDO Phase

ToUndo={lastLSNs of all Xacts in the Xact Table}
a.k.a. "losers"

Repeat:
- Choose (and remove) largest LSN among ToUndo.
- If this LSN is a CLR and undonextLSN==NULL
  - Write an End record for this Xact.
- If this LSN is a CLR, and undonextLSN != NULL
  - Add undonextLSN to ToUndo
- Else this LSN is an update. Undo the update, write a CLR, add prevLSN to ToUndo.

Until ToUndo is empty.

NOTE: This is simply a performance optimization on the naïve solution to do it in one backwards pass of the log!

Example of Recovery

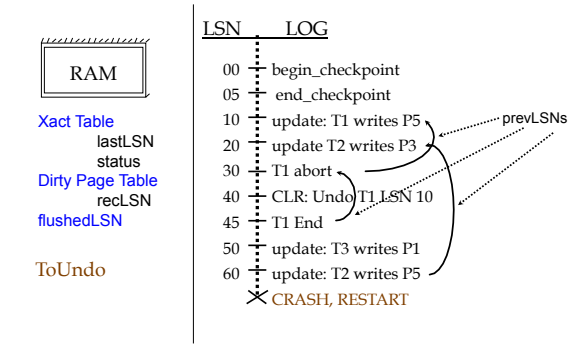

Example: Crash During Restart!

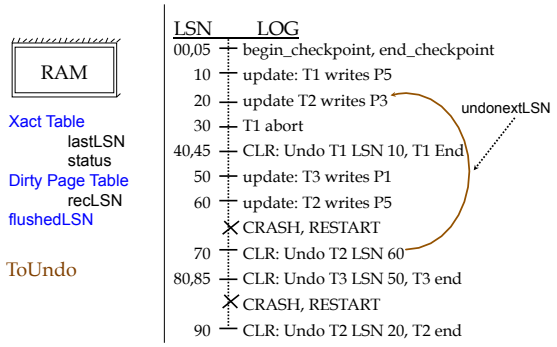

Additional Crash Issues

- **What happens if system crashes during Analysis? During REDO?**
- **How do you limit the amount of work in REDO?**
  - Flush asynchronously in the background.
  - Watch "hot spots"!
- **How do you limit the amount of work in UNDO?**
  - Avoid long-running Xacts.

Summary of Logging/Recovery

- **Recovery Manager guarantees Atomicity & Durability.**
- **Use WAL to allow STEAL/NO-FORCE w/o sacrificing correctness.**
- **LSNs identify log records; linked into backwards chains per transaction (via prevLSN).**
- **pageLSN allows comparison of data page and log records.**

Summary, Cont.

- **Checkpointing: A quick way to limit the amount of log to scan on recovery.**
- **Recovery works in 3 phases:**
  - Analysis: Forward from checkpoint.
  - Redo: Forward from oldest recLSN.
  - Undo: Backward from end to first LSN of oldest Xact alive at crash.
- **Upon Undo, write CLRs.**
- **Redo "repeats history": Simplifies the logic!**